

Subodh Ashok Bhyri

subodhbhyri811@gmail.com | (720) 666-5103 | www.linkedin.com/in/subodhbhyri | <https://github.com/subodhbhyri>

Portfolio - <https://subodhbhyri.vercel.app>

EDUCATION

University of Florida

Master of Science, Computer Science, GPA - 3.6/4

Aug 2024 - May 2026

Gainesville, Florida

University of Mumbai

Bachelor of Engineering, Computer Science, GPA - 8.54/10

Aug 2019 - Jul 2023

Mumbai, India

WORK EXPERIENCE

Vishwanath Projects Limited

Aug 2023 - Apr 2024

Associate Software Engineer

Hyderabad, India

- Architected and delivered a centralized admin screen for real-time monitoring of transmission line health, aggregating sensor data to improve operational visibility and reduce incident response time
- Designed and implemented a scalable RBAC system using React and Node.js, enabling secure user management for 500+ users across multi-tier enterprise environments
- Optimized data transfers between platform components via REST API integration using Node.js, achieving a 60% increase in data retrieval speed and lowering latency by 20ms
- Secured user authentication by integrating OAuth2 and JWT, resulting in a 99.9% success rate for secure token-based access and an improved user authentication workflow

OPEN-SOURCE CONTRIBUTIONS

- [chartjs/Chart.js](#) — Fixed missing [number, number] tuple type support for line/scatter datasets
- [visgl/deck.gl](#) — Added TypeScript types for loaders.gl point cloud data format (MeshAttributes)
- [visgl/react-map-gl](#) — Removed stale @ts-ignore directives for localIdeographFontFamily across all map modules

PROJECTS

LLM Eval: Multi-Model AI Evaluation Platform | Python, FastAPI, React/TypeScript, PostgreSQL, Redis, Kafka, RAGAS, Docker, Prometheus | <https://ragas-bench.vercel.app>

- Built a high-throughput evaluation pipeline processing concurrent LLM requests via async Python, designing fault-tolerant storage layers using PostgreSQL and Redis with sub-100ms cache hit latency and automatic failover
- Shipped a React/TypeScript dashboard with metric cards, radar chart, composite bar chart, and 5-second real-time polling with status badge updates, deployed on Vercel with FastAPI on Render
- Instrumented infrastructure with Prometheus metrics at /metrics and integrated Kafka for real-time event streaming, enabling proactive monitoring and capacity planning across distributed evaluation services

Timeline: Explore History Through Time and Space | React Globe GL, MapLibre GL JS, Gemini API, MongoDB Atlas Vector Search (RAG), MCP (Model Context Protocol) Server | <https://timeline-tau-eight.vercel.app>

- Architected scalable storage and retrieval infrastructure using MongoDB Atlas Vector Search and MCP, designing fault-tolerant data access layers with tiered fallback logic ensuring high availability for downstream AI agents
- Implemented a Semantic RAG pipeline using Gemini embeddings and Atlas Vector Search with tiered fallback logic, reducing LLM hallucinations by 40%

Comrades: Vicinity-Based Social Networking Platform | Node.js, Express.js, Mongoose, JWT, Socket.io, Google Maps API Geospatial Queries, LangChain, Gemini | <https://comrades.vercel.app>

- Built a location-based platform with scalable storage layers using MongoDB geospatial indexes and Node.js/Express, supporting high-throughput event processing across 1,000+ simulated concurrent user sessions
- Built an AI agent using LangChain and Gemini to classify post intent, query nearby users, and dispatch targeted real-time notifications via Socket.io, achieving 95%+ classification accuracy with sub-second latency

Slotify: AI-Powered Distributed Ticket Reservation System | Java 21, Spring Boot, Apache Kafka, Redis, PostgreSQL, WebSockets, gRPC, Kubernetes, Prometheus, Grafana, GitHub Actions | <https://github.com/subodhbhyri/slotify>

- Engineered scalable storage and caching infrastructure using Redis distributed locking and PostgreSQL for transactional consistency, handling high-contention concurrent booking requests with automatic seat reassignment
- Validated sub-35ms p99 booking latency under 100 concurrent users via k6 load testing, processing 65+ reservations/second with 100% success rate across 19,745 requests

SKILLS

Languages & Frameworks: Java (21), Python, TypeScript, JavaScript, Spring Boot, FastAPI, Node.js, gRPC, Scala

Storage & Infrastructure: Distributed Storage, Redis (Caching, Sorted Sets, Distributed Locking), PostgreSQL, MongoDB Atlas, Apache Cassandra, HBase, HDFS, Object Storage, SQL Server

Big Data & Streaming: Apache Kafka, PySpark, Hadoop, Hive, Apache Flink, Airflow, ETL/ELT, Real-time Data Pipelines

Cloud & DevOps: AWS, Supabase, Docker, Kubernetes, GitHub Actions (CI/CD), Prometheus, Grafana, OpenTelemetry

Backend & Architecture: Distributed Systems, High-Availability Design, Fault Tolerance, Concurrency & Multithreading