

# Subodh Ashok Bhyri

subodhbhyri811@gmail.com | (720) 666-5103 | [www.linkedin.com/in/subodhbhyri](https://www.linkedin.com/in/subodhbhyri) | <https://github.com/subodhbhyri>

Portfolio - <https://subodhbhyri.vercel.app>

## EDUCATION

### University of Florida

Master of Science, Computer Science, GPA - 3.6/4

Aug 2024 - May 2026

Gainesville, Florida

### University of Mumbai

Bachelor of Engineering, Computer Science, GPA - 8.54/10

Aug 2019 - Jul 2023

Mumbai, India

## WORK EXPERIENCE

### Vishwanath Projects Limited

Aug 2023 - Apr 2024

Associate Software Engineer

Hyderabad, India

- Architected and delivered a centralized admin screen for real-time monitoring of transmission line health, aggregating sensor data to improve operational visibility and reduce incident response time
- Designed and implemented a scalable RBAC system using React and Node.js, enabling secure user management for 500+ users across multi-tier enterprise environments
- Optimized data transfers between platform components via REST API integration using Node.js, achieving a 60% increase in data retrieval speed and lowering latency by 20ms
- Secured user authentication by integrating OAuth2 and JWT, resulting in a 99.9% success rate for secure token-based access and an improved user authentication workflow

## OPEN-SOURCE CONTRIBUTIONS

- [chartjs/Chart.js](#) — Fixed missing [number, number] tuple type support for line/scatter datasets
- [visgl/deck.gl](#) — Added TypeScript types for loaders.gl point cloud data format (MeshAttributes)
- [visgl/react-map-gl](#) — Removed stale @ts-ignore directives for localIdeographFontFamily across all map modules

## PROJECTS

**LLM Eval: Multi-Model AI Evaluation Platform** | Python, FastAPI, React/TypeScript, PostgreSQL, Redis, Kafka, RAGAS, Docker, Prometheus | <https://ragas-bench.vercel.app>

- Built an automated evaluation pipeline benchmarking Claude, GPT-4o, and Gemini concurrently, implementing RAGAS test suites scoring faithfulness, context precision, context recall, and answer relevancy across model versions
- Designed regression tracking infrastructure using PostgreSQL to persist evaluation runs and Redis caching for low-latency queries, enabling automated detection of model performance degradation across releases
- Instrumented backend with Prometheus metrics and integrated Kafka producer for real-time streaming of test results, enabling proactive alerting and CI/CD pipeline integration for automated eval runs

**Timeline: Explore History Through Time and Space** | React Globe GL, MapLibre GL JS, Gemini API, MongoDB Atlas Vector Search (RAG), MCP (Model Context Protocol) Server | <https://timeline-tau-eight.vercel.app>

- Built and validated an AI-native platform integrating MCP tool-calling and RAG pipelines, authoring integration test suites to verify agent behavior, retrieval accuracy, and end-to-end response correctness
- Implemented a Semantic RAG pipeline using Gemini embeddings and Atlas Vector Search with tiered fallback logic, reducing LLM hallucinations by 40%

**Comrades: Vicinity-Based Social Networking Platform** | Node.js, Express.js, Mongoose, JWT, Socket.io, Google Maps API Geospatial Queries, LangChain, Gemini | <https://comrades.vercel.app>

- Developed a location-based social platform using MongoDB geospatial queries and Node.js/Express, authoring integration test suites to validate real-time Socket.io interactions across 1,000+ simulated concurrent sessions
- Built an AI agent using LangChain and Gemini to classify post intent, query nearby users, and dispatch targeted real-time notifications via Socket.io, achieving 95%+ classification accuracy with sub-second latency

**Slotify: AI-Powered Distributed Ticket Reservation System** | Java 21, Spring Boot, Apache Kafka, Redis, PostgreSQL, WebSockets, gRPC, Kubernetes, Prometheus, Grafana, GitHub Actions | <https://github.com/subodhbhyri/slotify>

- Engineered an event-driven ticket reservation platform using Java 21 and Kafka, simulating 1M+ concurrent booking attempts to perform stress and regression testing across distributed microservices
- Validated sub-35ms p99 booking latency under 100 concurrent users via k6 load testing, processing 65+ reservations/second with 100% success rate across 19,745 requests

## SKILLS

**Languages & Frameworks:** Java (21), Python, TypeScript, JavaScript, Spring Boot, FastAPI, Node.js, Express, React, gRPC

**Testing & Quality:** RAGAS, JUnit, Jest, TDD, BDD, Unit Testing, Integration Testing, Regression Testing, Load Testing (k6, JMH), End-to-End Testing, Test Data Pipelines, Code Review

**Automation & CI/CD:** GitHub Actions, Docker, Kubernetes, Airflow, OpenTelemetry, CI/CD Pipelines, Automated Alerting

**Databases & Cloud:** AWS, Supabase, Docker, Kubernetes, PostgreSQL, MySQL, MongoDB, Redis, Apache Cassandra

**Backend & Architecture:** Distributed Systems, Microservices, Event-Driven Architecture (Kafka), REST APIs, OAuth2/JWT